First 15 minutes

Read rules, objective, input/output and official example. Confirm Python version, AI/prewritten-code eligibility, data rights, dependencies and resource limits. Run the untouched official baseline. Do not substitute past contest assumptions.

Market and execution

Spread = ask - bid. Midpoint is a mark, not an executable fill. Track cash + inventory * mark, fees and pending-order exposure. A cancel request does not free risk until acknowledged. Test duplicate and partial fills. Compare no-op, fixed and inventory-aware quotes. Synthetic fills prove no edge.

Tabular ML

Configure target, features and time explicitly. Sort in time; keep train/validation/test separate. Fit imputation and scaling only on train. Purge overlapping labels and embargo gaps. Compare constant with ridge/logistic. Select settings on validation; leave test untouched. Predictive accuracy is not net trading profit.

Allocation and games

Write legal choices and the exact payoff first. Use enumeration as a tiny-case oracle, then knapsack DP for integer costs. Shared rewards depend on rival counts. Auctions need explicit payment and tie rules. Zero-sum games need maximin bounds; general-sum games require different analysis.

Coding and packaging

Start with an exact small oracle. Dijkstra requires nonnegative edges; Bellman-Ford detects reachable negative cycles. Confirm filename, function/class signature and import restrictions. Static AST checks do not prove runtime or sandbox safety. Run official samples and edge cases under the real limits.

Experiment and handoff

Change one idea at a time. Save config, seed, input hash, source hash and metric. Compare identical cases, then new cases and pessimistic assumptions. Preserve last-known-good submission. Before stopping: run arena validate, save results, generate handoff, list three next actions. External AI memory remains disabled.